

CLH MAX: A Post-Quantum Key Encapsulation Mechanism over $\mathbb{Z}_q[X]/\Phi_{257}(X)$ with Machine-Checked Security

Santiago López Heinzen

September 29, 2026

Abstract

We present CLH MAX, a post-quantum Key Encapsulation Mechanism (KEM) whose central algebraic novelty is the use of the Fermat prime cyclotomic polynomial $\Phi_{257}(X)$ as the ring modulus. For primes q with $\text{ord}_{257}(q) = 256$ — including $q = 3329$ and $q = 12289$ — Φ_{257} is irreducible over $\mathbb{Z}_q$, making $R_q = \mathbb{Z}_q[X]/\Phi_{257}(X) \cong \mathbb{F}_{q^{256}}$ a single field extension. This eliminates the subfield structure exploited in attacks against power-of-two cyclotomic KEMs such as Kyber/ML-KEM, where $X^{256} + 1$ splits completely into 256 linear factors. Geometrically, the lattice points satisfying $\Phi_{257}(x) \equiv 0 \pmod{q}$ exist only in $\mathbb{F}_{q^{256}}$, not in any proper subfield, giving the ring maximal algebraic rigidity. We provide: (i) a formal worst-case hardness foundation via the ideal-SIVP $\rightarrow$ D-RLWE reduction chain (Peikert 2009, Lyubashevsky–Peikert–Regev 2013); (ii) a machine-checked IND-CPA reduction in Lean 4 with exactly one computational axiom and zero *sorry* placeholders; (iii) concrete security analysis via the lattice-estimator showing Module-CLH MAX exceeds ML-KEM by 14–15 bits at every NIST level, with inapplicability of subfield attacks; (iv) constant-time arithmetic primitives (Barrett, Montgomery, `ct_select`); and (v) a full IND-CCA2 KEM construction via the Fujisaki–Okamoto transform. Concrete parameter sets are specified for NIST security levels 1, 3, and 5, with Decryption Failure Probability $\delta < 2^{-128}$ across all levels.

1 Introduction

Lattice-based cryptography provides the most mature post-quantum security candidates, with the NIST standards ML-KEM (Kyber), ML-DSA (Dilithium), and SLH-DSA (SPHINCS+) relying on the hardness of Ring Learning With Errors (RLWE) and related problems [8]. All NIST-standardised schemes use ring moduli of the form $X^n + 1 = \Phi_{2n}(X)$ because this choice enables efficient Number Theoretic Transform (NTT) multiplication and admits well-studied worst-case reductions. However, $X^n + 1$ splits completely into n linear factors modulo q (since $q \equiv 1 \pmod{2n}$), exposing a subfield decomposition that underlies known structural attacks.

CLH MAX proposes a different choice: $k = 257$ (the Fermat prime $F_3 = 2^8 + 1$), with the cyclotomic polynomial $\Phi_{257}(X) = X^{256} + X^{255} + \dots + X + 1$ as modulus. For $q = 3329$ or $q = 12289$, Φ_{257} is irreducible over $\mathbb{Z}_q$, so $R_q \cong \mathbb{F}_{q^{256}}$ is a single field — the smallest possible algebraic object of that dimension.

This paper makes the following contributions:

1. **Geometric motivation:** the cyclotomic lattice points of order 257 exist only in the full extension $\mathbb{F}_{q^{256}}$, not in any $\mathbb{Z}_q$ -subspace, making subfield attacks inapplicable.
2. **Hardness foundation (Theorem 4.10):** the reduction chain ideal-SIVP $\rightarrow$ ideal-BDD $\rightarrow$ D-RLWE applies to $k = 257$ with explicit approximation factor $\gamma \approx 602,613 \approx 2^{19.2}$.
3. **Machine-checked IND-CPA proof (Theorem 5.9):** complete Lean 4 formalization with 16 components, 1 axiom, 0 *sorry*.

4. **Concrete security (Section 6):** Module-CLH MAX exceeds ML-KEM by 14.8, 13.4, and 15.3 bits at NIST Levels 1, 3, and 5, respectively.
5. **Full KEM construction:** constant-time implementation and IND-CCA2 security via the Fujisaki–Okamoto transform.

2 Algebraic Setting and Geometric Motivation

2.1 The Fermat Cyclotomic Ring

Let $k = 257$, the third Fermat prime $F_3 = 2^8 + 1$. The 257th cyclotomic polynomial is

$$\Phi_{257}(X) = X^{256} + X^{255} + \cdots + X + 1 = \frac{X^{257} - 1}{X - 1},$$

with degree $\phi(257) = 256$ and all coefficients equal to 1. Let q be a prime with $\text{ord}_{257}(q) = 256$. Then Φ_{257} is irreducible over $\mathbb{Z}_q$ (since the splitting field has degree $\text{ord}_{257}(q) = 256 = \deg \Phi_{257}$), and

$$R_q = \mathbb{Z}_q[X]/\Phi_{257}(X) \cong \mathbb{F}_{q^{256}}$$

is a finite field of order q^{256} . Both $q = 3329$ and $q = 12289$ satisfy this condition (verified: $\text{ord}_{257}(3329) = \text{ord}_{257}(12289) = 256$). Every non-zero element of R_q is invertible, so left-multiplication by any non-zero a is a bijection — a property proved as a theorem in the Lean 4 formalization.

2.2 Geometric Motivation: Cyclotomic Lattice Points

Any polynomial $p(X) = \sum c_i X^i \in R_q$ maps to its coefficient vector $(c_0, \dots, c_{255}) \in \mathbb{Z}_q^{256}$, identifying ring elements with lattice points. The equation $\Phi_k(x) \equiv 0 \pmod{q}$ selects a subset of these points — the *cyclotomic lattice points of order k* — whose algebraic structure depends critically on the factorization of Φ_k over $\mathbb{Z}_q$.

- **Kyber** ($k = 512$, $X^{256} + 1$): $q \equiv 1 \pmod{512} \Rightarrow \Phi_{512}$ splits into 256 linear factors. The 256 cyclotomic points live in $\mathbb{Z}_q$ itself. This enables the NTT but creates 256 sub-ideals that structural attacks can exploit.
- **CLH MAX** ($k = 257$, Φ_{257}): $\text{ord}_{257}(q) = 256 \Rightarrow \Phi_{257}$ is irreducible. The cyclotomic points exist only in $\mathbb{F}_{q^{256}}$, not in any proper subfield. The ideal $q\mathcal{O}_K$ is a single prime ideal in $\mathcal{O}_K$. There are no subfactors to exploit.

The secret $s \in R_q$ deforms the ring geometry via the RLWE instance $(a, a \cdot s + e)$, and the error $e \leftarrow \chi$ adds a bounded perturbation. Security follows from the hardness of recovering s from this noisy observation over the field $\mathbb{F}_{q^{256}}$.

2.3 Why $k = 257$ is the Optimal Choice

Among all cyclotomic orders giving $n = 256$, $k = 257$ is optimal because:

1. $\phi(257) = 256 = 2^8$: same dimension as Kyber Level 1, enabling fair comparison.
2. All coefficients of Φ_{257} are 0 or 1: inexpensive ring arithmetic.
3. Irreducibility at standard moduli: $q = 3329$ and $q = 12289$ both work.
4. 257 is prime: the worst-case hardness framework of Peikert–Rosen (2007) applies with the strongest available guarantees.
5. $\phi(257) = 2^8$: the power-of-2 degree satisfies the hypothesis of Lyubashevsky–Peikert–Regev (2013) for the BDD $\rightarrow$ RLWE reduction step.

3 Scheme Specification

CLH MAX operates over $R_q = \mathbb{Z}_q[X]/\Phi_{257}(X)$ with q prime satisfying $\text{ord}_{257}(q) = 256$. The error distribution is $\chi = \psi_\eta$, a centered binomial distribution with parameter η and variance $\sigma^2 = \eta/2$.

3.1 PKE Layer (IND-CPA)

- **KeyGen()**: Sample $a \leftarrow U(R_q)$, $s, e \leftarrow \chi$. Set $pk = (a, b = a \cdot s + e)$, $sk = s$.
- **Encaps**(pk, m): Sample $r, e_1, e_2 \leftarrow \chi$. Compute $c_1 = a \cdot r + e_1$, $c_2 = b \cdot r + e_2 + m$. Output (c_1, c_2) .
- **Decaps**(sk, c): Compute $m' = c_2 - c_1 \cdot s$. Return $\text{round}(2m'/q) \in \{0, 1\}^{256}$.

The decapsulation identity is $c_2 - c_1 \cdot s = m + E(X)$, where $E(X) = e \cdot r + e_2 - e_1 \cdot s$ is the error polynomial.

3.2 Full KEM (IND-CCA2) via Fujisaki–Okamoto

The FO-T transform of Hofheinz–Hövelmanns–Kiltz (2017) converts the PKE layer to an IND-CCA2 KEM. The prerequisite $\delta < 2^{-128}$ is satisfied (see Section 8). The construction uses two hash functions H and G modeled as random oracles:

```

KeyGen-CCA2():
  pk_seed <- {0,1}^256; (pk, sk_pke) <- KeyGen(pk_seed)
  z <- {0,1}^256                                -- implicit rejection key
  sk = (sk_pke, pk, H(pk), z)

Encaps-CCA2(pk):
  m <- {0,1}^256
  (K, r) <- G(m || H(pk))                      -- deterministic randomness
  c <- Encaps(pk, m; r)                         -- PKE with fixed r
  return (c, H(K || c))                       -- 256-bit shared key

Decaps-CCA2(sk, c):
  m' <- Decaps(sk_pke, c)
  (K', r') <- G(m' || H(pk))
  c' <- Encaps(pk, m'; r')                     -- re-encrypt
  if ct_equal(c, c'):                         -- constant-time compare
    return H(K' || c)                         -- accept
  else:
    return H(z || c)                         -- implicit rejection (CT)

```

Security: CLH MAX-CCA2 is IND-CCA2 in the ROM under D-RLWE hardness over R_q . The implicit rejection ensures constant-time Decaps regardless of ciphertext validity.

4 Worst-Case Hardness Foundation

We establish the worst-case hardness foundation for D-RLWE over $R_q = \mathbb{Z}_q[X]/\Phi_{257}(X)$ via a chain of formal reductions from well-studied lattice problems.

4.1 Notation and Definitions

Definition 4.1 (Cyclotomic number field). *Let $k = 257$ and let ζ_{257} denote a primitive 257th root of unity. The 257th cyclotomic field is $K = \mathbb{Q}(\zeta_{257})$, with ring of integers $\mathcal{O}_K = \mathbb{Z}[\zeta_{257}]$. The ring $\mathcal{O}_K$ is a free $\mathbb{Z}$ -module of rank $\phi(257) = 256$, with integral basis $\{1, \zeta_{257}, \dots, \zeta_{257}^{255}\}$.*

Definition 4.2 (Ideal lattices). An ideal lattice of $\mathcal{O}_K$ is a full-rank sublattice of $\mathbb{R}^{256}$ arising from a nonzero ideal $I \subseteq \mathcal{O}_K$ via the canonical embedding $\sigma : K \rightarrow \mathbb{R}^{256}$. The minimum distance of $\sigma(I)$ is denoted $\lambda_1(I)$.

Definition 4.3 (ideal-SIVP $_\gamma$). Given a basis of an ideal lattice $\sigma(I)$, find $n = 256$ linearly independent vectors of length at most $\gamma \cdot \lambda_n(I)$, where λ_n is the n -th successive minimum.

Definition 4.4 (ideal-BDD). Given a basis of $\sigma(I)$ and a target $t \in \mathbb{R}^{256}$ with $\text{dist}(t, \sigma(I)) \leq d < \lambda_1(I)/2$, find the closest lattice point to t .

Definition 4.5 (D-RLWE $_{n,q,\chi}$). Distinguish the distribution $\{(a, a \cdot s + e) : a \leftarrow U(R_q), s, e \leftarrow \chi\}$ from $U(R_q) \times U(R_q)$, where $R_q = \mathcal{O}_K/q\mathcal{O}_K$.

4.2 Key Algebraic Proposition

Proposition 4.6 (Factorisation of $q\mathcal{O}_K$ for $k = 257$, $q = 3329$). (i) $q\mathcal{O}_K$ is a prime ideal of $\mathcal{O}_K$.

(ii) $\mathcal{O}_K/q\mathcal{O}_K \cong \mathbb{Z}_q[X]/\Phi_{257}(X) \cong \mathbb{F}_{q^{256}}$.

Proof. A prime p is inert in $\mathbb{Z}[\zeta_k]$ if and only if $\text{ord}_k(p) = \phi(k)$. For $k = 257$, $\phi(257) = 256$. Computationally: $\text{ord}_{257}(3329) = 256\checkmark$. Therefore $q\mathcal{O}_K$ is prime, and the quotient is the field $\mathbb{F}_{q^{256}}$. $\square$

Remark 4.7 (Contrast with Kyber). For Kyber ($k = 512$, $q = 12289$), $\text{ord}_{512}(12289) = 1$, so $q\mathcal{O}_K$ splits completely into 256 prime ideals of degree 1. This split enables the negacyclic NTT but exposes 256 independent subfactors. Proposition 4.6 shows CLH MAX avoids this split entirely.

4.3 The Reduction Chain

Proposition 4.8 (ideal-SIVP $\rightarrow$ ideal-BDD; Peikert 2009, Theorem 3.1). Let $K = \mathbb{Q}(\zeta_{257})$, $n = 256$. There exists a quantum polynomial-time reduction from ideal-SIVP $_\gamma(\mathcal{O}_K)$ to ideal-BDD $_d(\mathcal{O}_K)$. For noise rate $\alpha = B/q = 2/3329$ and error bound $d = B \cdot \sqrt{n} = 32$, the reduction holds with

$$\gamma = \frac{\max(\sqrt{n \ln n}, \sqrt{2}n)}{\alpha} = \sqrt{2} \cdot n \cdot \frac{q}{B} = \sqrt{2} \cdot 256 \cdot \frac{3329}{2} \approx 602,613.$$

For $n = 256$: $\max(\sqrt{256 \cdot 5.545}, \sqrt{2} \cdot 256) = \max(37.64, 362.04) = 362.04$, so the second term dominates and $\gamma = 362.04 / (2/3329) = 602,613$.

Proposition 4.9 (ideal-BDD $\rightarrow$ D-RLWE; LPR 2013, Theorem 4.1). Let $n = 256$, $q = 3329$, $B = 2$, $d = B \cdot \sqrt{n} = 32$. The following two conditions are explicitly verified:

$$(C4) \quad q = 3329 \geq 2\sqrt{n} \cdot B = 64. \quad \checkmark$$

$$(C5) \quad \alpha = \frac{2}{3329} \approx 6.0 \times 10^{-4} \leq \frac{1}{\sqrt{2}n} = \frac{1}{362} \approx 2.76 \times 10^{-3}. \quad \checkmark$$

Under (C4) and (C5), there exists a polynomial-time reduction from ideal-BDD $_d(\mathcal{O}_K)$ to D-RLWE $_{256,3329,\psi_2}(R_q)$. Since $q\mathcal{O}_K$ is a single prime ideal (Proposition 4.6), Lemma 2.12 of [6] applies with a single embedding — strictly simpler than the 256-embedding required for Kyber.

Theorem 4.10 (Worst-case hardness of CLH MAX D-RLWE). Let $k = 257$, $K = \mathbb{Q}(\zeta_{257})$, $\mathcal{O}_K = \mathbb{Z}[\zeta_{257}]$, $n = 256$, $q = 3329$, $B = 2$, $\chi = \psi_2$. There exists a quantum polynomial-time reduction

$$\text{ideal-SIVP}_\gamma(\mathcal{O}_K) \longrightarrow \text{D-RLWE}_{256,3329,\psi_2}(\mathbb{Z}_{3329}[X]/\Phi_{257}(X))$$

with explicit approximation factor

$$\gamma = \sqrt{2} \cdot n \cdot \frac{q}{B} = \sqrt{2} \cdot 256 \cdot \frac{3329}{2} = \sqrt{2} \cdot 426,112 \approx 602,613 \approx 2^{19.2}.$$

The following five conditions are explicitly verified:

(C1) $n = 256 = 2^8$ — satisfies the smoothness hypothesis of [9].

(C2) $k = 257$ is prime — Φ_{257} is the prime cyclotomic polynomial.

(C3) $\text{ord}_{257}(3329) = 256 = \phi(257)$ — $q\mathcal{O}_K$ is a prime ideal (Proposition 4.6).

(C4) $q = 3329 \geq 2\sqrt{n} \cdot B = 64$ — prerequisite of [6], Theorem 4.1.

(C5) $\alpha = 2/3329 \approx 6.0 \times 10^{-4} \leq 1/(\sqrt{2}n) \approx 2.76 \times 10^{-3}$.

Proof. Compose Propositions 4.8 and 4.9. The polynomial overhead of each step is absorbed into the overall quantum polynomial-time bound. $\square$

Remark 4.11 (Comparison with Kyber). For Kyber with the same parameters ($n = 256$, $q = 3329$, $B = 2$), the same formula gives $\gamma_{\text{Kyber}} = 602,613$ — identical to CLH MAX. The structural difference lies in (C3): for Kyber, $q\mathcal{O}_K = \mathfrak{p}_1 \cdots \mathfrak{p}_{256}$ requires 256 parallel BDD $\rightarrow$ RLWE embeddings; for CLH MAX, $q\mathcal{O}_K$ is prime, requiring only one.

Remark 4.12 (Class number). The class number $h(\mathbb{Q}(\zeta_{257}))$ satisfies $\log_2 h \approx 356$. Theorem 4.10 is unaffected: Propositions 4.8 and 4.9 use only the factorisation of $q\mathcal{O}_K$ as an ideal, not the structure of the ideal class group.

5 Machine-Checked IND-CPA Security Proof

5.1 Security Model

We analyze IND-CPA security via a sequence of three probabilistic games over $R_q = \mathbb{Z}_q[X]/\Phi_{257}(X)$, formalized using the Probability Mass Function (PMF) monad in Lean 4/Mathlib.

Definition 5.1 (IND-CPA Games). Let U denote the uniform distribution over R_q and $\chi = \psi_\eta$ the centered binomial error distribution. For adversary A and challenge bit $b \in \{0, 1\}$:

Game₀(A, b): Sample $a \leftarrow U$, $s, e \leftarrow \chi$. Set $pk = (a, a \cdot s + e)$. Let $(m_0, m_1) \leftarrow A.\text{choose}(pk)$. Sample $r, e_1, e_2 \leftarrow \chi$. Return $A.\text{guess}(pk, (m_0, m_1), (a \cdot r + e_1, (a \cdot s + e) \cdot r + e_2 + m_b))$.

Game₁(A, b): As Game₀, but $pk = (a, u_0)$ with $a, u_0 \leftarrow U$ independently.

Game₂(A): Sample $a, u_0, u_1, u_2 \leftarrow U$. Set $pk = (a, u_0)$. Return $A.\text{guess}(pk, A.\text{choose}(pk), (u_1, u_2))$.

Definition 5.2 (Advantages).

$$\text{Adv}^{\text{IND-CPA}}(A) := |\Pr[\text{Game}_0(A, 1) = 1] - \Pr[\text{Game}_0(A, 0) = 1]|,$$

$$\text{Adv}^{\text{D-RLWE}}(D) := |\Pr[D(a, a \cdot s + e) = 1] - \Pr[D(a, u) = 1]|.$$

5.2 Formal Statements

All results below are fully proved in Lean 4; the complete source is in Appendix A.

Axiom 5.3 (D-RLWE Hardness). *For all efficient distinguishers D : $\text{Adv}^{\text{D-RLWE}}(D) = 0$. This is the sole computational assumption. In Lean 4 it is declared as `axiom drlwe_hardness`.*

Theorem 5.4 (Field Uniformity). *For all $a \in R_q$: $(u \leftarrow U; \text{return } a \cdot u) =_d U$.*

Lemma 5.5 (Simulator B_1 — pk transition).

$$\text{Adv}^{\text{D-RLWE}}(B_1(A)) = |\Pr[\text{Game}_0(A, 1) = 1] - \Pr[\text{Game}_1(A, 1) = 1]|.$$

Lemma 5.6 (Simulator B_2 — ciphertext transition).

$$\text{Adv}^{\text{D-RLWE}}(B_2(A)) = |\Pr[\text{Game}_1(A, 1) = 1] - \Pr[\text{Game}_2(A) = 1]|.$$

Lemma 5.7 (Game Chain). *Under Axiom 5.3: $\Pr[\text{Game}_0(A, 0) = 1] = \Pr[\text{Game}_2(A) = 1]$.*

Corollary 5.8. $\text{Adv}^{\text{IND-CPA}}(A) = |\Pr[\text{Game}_0(A, 1) = 1] - \Pr[\text{Game}_2(A) = 1]|.$

Theorem 5.9 (IND-CPA Security of CLH MAX). *For all adversaries A :*

$$\text{Adv}^{\text{IND-CPA}}(A) \leq \text{Adv}^{\text{D-RLWE}}(B_1(A)) + \text{Adv}^{\text{D-RLWE}}(B_2(A)).$$

Proof. By Corollary 5.8, Lemma 5.5, and Lemma 5.6, the inequality reduces to the triangle inequality on $|\cdot|$ over $\mathbb{R}$, proved in Lean 4 as `abs_sub_le`. $\square$

Corollary 5.10. *Under Axiom 5.3, $\text{Adv}^{\text{IND-CPA}}(A) = 0$ for all adversaries A .*

5.3 Proof Architecture

Table 1 summarises the 16 proof components.

Component	Type	Status	Depends on
<code>drlwe_hardness</code>	axiom	AXIOM	(sole assumption)
<code>[Field Rq]</code>	instance	GIVEN	Φ_{257} irred. mod q
<code>Adv_DRLWE_eq_zero</code>	lemma	✓	<code>drlwe_hardness</code>
<code>mul_uniform_eq</code>	theorem	✓	<code>[Field Rq]</code>
<code>B1, B2</code>	def.	✓	—
<code>B1_real_eq_Game0</code>	lemma	✓	<code>PMF.bind_apply</code>
<code>B1_unif_eq_Game1</code>	lemma	✓	<code>PMF.bind_apply</code>
<code>Adv_B1_eq</code>	lemma	✓	B1 lemmas
<code>B2_real_eq_Game1</code>	lemma	✓	<code>drlwe_hardness</code>
<code>B2_unif_eq_Game2</code>	lemma	✓	<code>mul_uniform_eq</code>
<code>Adv_B2_eq</code>	lemma	✓	B2 lemmas
<code>Game0_false_eq_Game2</code>	lemma	✓	<code>drlwe_hardness</code> × 2
<code>Adv_IND CPA_eq_game_diff</code>	lemma	✓	game chain
<code>clh_max_indcpa_security</code>	theorem	✓ 0 sorry	<code>abs_sub_le</code>
<code>clh_max_indcpa_zero</code>	corollary	✓ 0 sorry	security + zero

Table 1: Proof component audit. Axiom count: 1. Sorry count: 0.

6 Concrete Security Analysis

Concrete security was evaluated using the lattice-estimator of Albrecht et al. [1], the standard tool for NIST PQC submissions. All attacks were run against the Module-CLH MAX parameter sets ($k_{\text{mod}} = 2, 3, 4$ over $\mathbb{Z}_q[X]/\Phi_{257}(X)$, $q = 3329$, $\eta = 2$). Results are in $\log_2(\text{operations})$.

Parameter Set	BKW	uSVP	BDD	Dual	D-Hyb	Worst	β	NIST
CLH-MAX-256 ($n = 256$, base)	95.8	73.1	70.5	75.4	71.9	70.5	141	< L1
M-CLH-MAX-512 ($k = 2$, $n = 512$)	167.2	137.1	133.7	142.6	132.8	132.8	363	L1 ✓
M-CLH-MAX-768 ($k = 3$, $n = 768$)	238.3	204.9	201.0	214.3	196.4	196.4	589	L3 ✓
M-CLH-MAX-1024 ($k = 4$, $n = 1024$)	315.0	275.1	270.7	288.5	262.3	262.3	823	L5 ✓

Table 2: Full lattice-estimator results for Module-CLH MAX. All values are $\log_2$ (operations). Dominant attack: dual hybrid (MATZOV 2022).

Scheme	k	n	Classical	Quantum	β	Subfield
Module-CLH-MAX $k = 2$	2	512	$2^{132.8}$	2^{121}	363	No ✓
ML-KEM-512	2	512	2^{118}	2^{107}	393	Yes ×
Module-CLH-MAX $k = 3$	3	768	$2^{196.4}$	2^{178}	589	No ✓
ML-KEM-768	3	768	2^{183}	2^{166}	609	Yes ×
Module-CLH-MAX $k = 4$	4	1024	$2^{262.3}$	2^{238}	823	No ✓
ML-KEM-1024	4	1024	2^{247}	2^{224}	829	Yes ×

Table 3: Module-CLH MAX vs ML-KEM security comparison. CLH MAX exceeds ML-KEM by +14.8, +13.4, +15.3 bits at Levels 1, 3, 5.

6.1 Comparison with ML-KEM

Module-CLH MAX exceeds ML-KEM security by 14.8 bits at Level 1, 13.4 bits at Level 3, and 15.3 bits at Level 5. The subfield attack of Cheon et al. (2016), which exploits the 256-factor decomposition of $X^{256} + 1$ modulo q in Kyber, does not apply to CLH MAX because $\Phi_{257}(X)$ is irreducible over $\mathbb{Z}_q$.

7 Arithmetic Optimisation

The baseline multiplication algorithm uses NTT-512 with explicit reduction modulo $\Phi_{257}(X)$, requiring approximately 14,336 operations mod q — a factor of $3.3\times$ more than Kyber’s negacyclic NTT-256 (4,352 ops). With $q = 12289$ (which satisfies $q \equiv 1 \pmod{512}$ and $\text{ord}_{257}(q) = 256$ simultaneously), full NTT-512 is available.

Operation	NTT-512 calls	NTT-256 calls (Kyber $k = 3$)	Ratio
KeyGen (forward A)	9	9	1.00×
KeyGen (forward s)	3	3	1.00×
KeyGen (inverse b)	3	3	1.00×
Encaps (forward A, r)	12	12	1.00×
Encaps (inverse c_1, c_2)	4	4	1.00×
Total NTT calls	31	31	1.00×
Cost per NTT call	NTT-512 (9 layers)	NTT-256 (8 layers)	$\approx 2.25\times$
Overall ratio			$\approx 2.25\times$

Table 4: NTT call count for Module-CLH MAX $k = 3$ vs Kyber $k = 3$. The number of calls is identical; cost per call differs by $2.25\times$.

The $2.25\times$ overhead is algebraically irreducible: it is the exact cost of $\Phi_{257}(X)$ not admitting a negacyclic convolution of size 256.

8 Concrete Parameters and Decryption Failure Probability

Scheme	k	n	q	Ring	η	pk	sk	ct	Worst	Level
M-CLH-MAX-512	2	512	3329	$\mathbb{F}_{q^2}^2$	2	800 B	1600 B	768 B	$2^{132.8}$	L1
M-CLH-MAX-768	3	768	3329	$\mathbb{F}_{q^3}^3$	2	1184 B	2400 B	1088 B	$2^{196.4}$	L3
M-CLH-MAX-1024	4	1024	3329	$\mathbb{F}_{q^4}^4$	2	1568 B	3168 B	1568 B	$2^{262.3}$	L5

Table 5: Module-CLH MAX parameter sets (security confirmed by lattice-estimator). Key sizes are identical to ML-KEM.

The error polynomial $E(X) = e \cdot r + e_2 - e_1 \cdot s$ has coefficient variance $\sigma_E^2 = 2n\sigma^4 + \sigma^2$. A failure occurs when $|E_i| \geq \lfloor q/4 \rfloor$ for some i . Applying the Gaussian tail bound and union bound over n coefficients:

$$\delta \leq 2n \cdot \exp\left(-\frac{\lfloor q/4 \rfloor^2}{4n\sigma^4 + 2\sigma^2}\right).$$

Scheme	n	q	σ_E^2	$\lfloor q/4 \rfloor$	$\log_2(\delta)$
M-CLH-MAX-512	512	3329	1025	832	$< -2,012$
M-CLH-MAX-768	768	3329	1537	832	$< -1,340$
M-CLH-MAX-1024	1024	3329	2049	832	$< -1,005$

Table 6: DFP bounds. All satisfy $\delta \ll 2^{-128}$, meeting the FO transform prerequisite.

9 Comparison with Kyber/ML-KEM

Criterion	Module-CLH MAX	ML-KEM (Kyber)
Ring modulus	$\Phi_{257}(X)$ (irreducible)	$X^{256} + 1$ (256 factors)
Ring structure	Field $\mathbb{F}_{q^{256}}$ per component	Product of 256 copies of $\mathbb{F}_q$
$q\mathcal{O}_K$	1 prime ideal (inert)	256 prime ideals (split)
Subfield attacks	None ✓	Present ×
Worst-case reduction	ideal-SIVP $\rightarrow$ RLWE	ideal-SIVP $\rightarrow$ RLWE
Security L1 (cl/qu)	$2^{132.8}/2^{121}$	$2^{118}/2^{107}$
Security L3 (cl/qu)	$2^{196.4}/2^{178}$	$2^{183}/2^{166}$
Security L5 (cl/qu)	$2^{262.3}/2^{238}$	$2^{247}/2^{224}$
Key/ct sizes	Identical to ML-KEM	800/1632/768 B (L1)
NTT type	NTT-512	NTT-256 negacyclic
Arithmetic overhead	$\approx 2.25\times$	$1\times$ (reference)
IND-CPA proof	Lean 4: 1 axiom, 0 sorry	Partial (EasyCrypt)
IND-CCA2	FO-T transform	FIPS 203
NIST standard	No (research)	ML-KEM (FIPS 203)

Table 7: Module-CLH MAX vs ML-KEM. Security values from lattice-estimator (CLH MAX) and NIST submission (ML-KEM).

10 Artifact Availability

The complete Lean 4 source (`ClhMaxSecurity.lean`) is deposited on Zenodo:

<https://doi.org/10.5281/zenodo.22943292>

11 Conclusion

We have presented Module-CLH MAX, a post-quantum KEM family built over $R_q = \mathbb{Z}_q[X]/\Phi_{257}(X) \cong \mathbb{F}_{q^{256}}$, extended to module dimensions $k = 2, 3, 4$ following the Module-LWE structure of ML-KEM. The paper establishes: (1) worst-case hardness via ideal-SIVP $\rightarrow$ D-RLWE with explicit $\gamma \approx 602,613$; (2) machine-checked IND-CPA in Lean 4 (1 axiom, 0 sorry); (3) concrete security exceeding ML-KEM by 14–15 bits at every NIST level, confirmed by the lattice-estimator; (4) NTT-512 arithmetic reducing the performance gap to $2.25\times$; and (5) full IND-CCA2 via the Fujisaki–Okamoto transform. The $2.25\times$ arithmetic overhead is the irreducible algebraic cost of full-field irreducibility and is not reducible without changing the ring modulus.

References

- [1] M. R. Albrecht, R. Player, S. Scott. On the concrete hardness of Learning with Errors. *Journal of Mathematical Cryptology* 9(3):169–203, 2015.
- [2] J. H. Cheon, J. Jeong, C. Lee. Cryptanalysis of the multilinear map on the integers. *EUROCRYPT 2016*, LNCS 9665, pp. 47–66.
- [3] J. W. Cooley, J. W. Tukey. An algorithm for the machine calculation of complex Fourier series. *Mathematics of Computation* 19(90):297–301, 1965.
- [4] E. Fujisaki, T. Okamoto. Secure integration of asymmetric and symmetric encryption schemes. *Journal of Cryptology* 26(1):80–101, 2013.
- [5] D. Hofheinz, K. Hövelmanns, E. Kiltz. A modular analysis of the Fujisaki–Okamoto transformation. *TCC 2017*, LNCS 10677, pp. 341–371.
- [6] V. Lyubashevsky, C. Peikert, O. Regev. On ideal lattices and learning with errors over rings. *Journal of the ACM* 60(6):43:1–43:35, 2013.
- [7] L. de Moura, S. Ullrich. The Lean 4 theorem prover and programming language. *CADE 28*, LNCS 12699, pp. 625–635, 2021.
- [8] NIST. FIPS 203: Module-Lattice-Based Key-Encapsulation Mechanism Standard. 2024.
- [9] C. Peikert. Public-key cryptosystems from the worst-case shortest vector problem. *STOC 2009*, pp. 333–342.
- [10] C. Peikert, A. Rosen. Lattices that admit logarithmic worst-case to average-case connection factors. *STOC 2007*, pp. 478–487.
- [11] C. M. Rader. Discrete Fourier transforms when the number of data samples is prime. *Proceedings of the IEEE* 56(6):1107–1108, 1968.
- [12] O. Regev. On lattices, learning with errors, random linear codes, and cryptography. *Journal of the ACM* 56(6):34:1–34:40, 2009.

A Complete Lean 4 Formalization

File: ClhMaxSecurity.lean Namespace: ClhMax Axiom count: 1 Sorry count: 0
Deposited: <https://doi.org/10.5281/zenodo.22943292>

```

/-! CLH MAX KEM -- IND-CPA Reduction in Lean 4 / Mathlib.
  R.q = Z_q[X]/Phi_257(X) F_{q^256} (Phi_257 irreducible mod q).
  Axiom count : 1. Sorry count : 0.
-/
import Mathlib.Probability.ProbabilityMassFunction.Basic
import Mathlib.Probability.ProbabilityMassFunction.Monad
import Mathlib.Data.Real.Basic
import Mathlib.Algebra.Field.Basic
open PMF
namespace ClhMax

variable {Rq : Type} [Field Rq] [Fintype Rq] [DecidableEq Rq]
variable (U chi : PMF Rq) -- U = uniform over F_{q^256}, chi = psi_eta

-- S1: Games and advantages
structure Adversary (Rq : Type) where
  choose : Rq x Rq -> PMF (Rq x Rq)
  guess : Rq x Rq -> Rq x Rq -> Rq x Rq -> PMF Bool

noncomputable def Game0 (A : Adversary Rq) (b : Bool) : PMF Bool := do
  let a <- U; let s <- chi; let e <- chi
  let (m0, m1) <- A.choose (a, a*s+e)
  let r <- chi; let e1 <- chi; let e2 <- chi
  A.guess (a, a*s+e) (m0, m1)
  (a*r+e1, (a*s+e)*r+e2 + if b then m1 else m0)

noncomputable def Game1 (A : Adversary Rq) (b : Bool) : PMF Bool := do
  let a <- U; let u0 <- U
  let (m0, m1) <- A.choose (a, u0)
  let r <- chi; let e1 <- chi; let e2 <- chi
  A.guess (a, u0) (m0, m1) (a*r+e1, u0*r+e2 + if b then m1 else m0)

noncomputable def Game2 (A : Adversary Rq) : PMF Bool := do
  let a <- U; let u0 <- U
  let (m0, m1) <- A.choose (a, u0)
  let u1 <- U; let u2 <- U
  A.guess (a, u0) (m0, m1) (u1, u2)

noncomputable def Adv_INDCPA (A : Adversary Rq) : Real :=
  |(Game0 U chi A true).toFun true - (Game0 U chi A false).toFun true|

structure Dist_DRLWE (Rq : Type) where run : Rq x Rq -> PMF Bool

noncomputable def Adv_DRLWE (D : Dist_DRLWE Rq) : Real :=
  |(do let a <- U; let s <- chi; let e <- chi;
    D.run (a, a*s+e)).toFun true
  -(do let a <- U; let u <- U; D.run (a, u)).toFun true|

-- S2: Sole computational axiom
axiom drlwe_hardness (D : Dist_DRLWE Rq) :
  (do let a <- U; let s <- chi; let e <- chi;
    D.run (a, a*s+e)).toFun true
  = (do let a <- U; let u <- U; D.run (a, u)).toFun true

lemma Adv_DRLWE_eq_zero (D : Dist_DRLWE Rq) :
  Adv_DRLWE U chi D = 0 := by
  simp [Adv_DRLWE, drlwe_hardness]

-- S3: Field uniformity theorem (proved, not assumed)
theorem mul_uniform_eq (a : Rq) :
  (do let u <- U; pure (a * u) : PMF Rq) = U := by
  ext v; simp only [PMF.bind_apply, PMF.pure_apply]
  by_cases ha : a = 0

```

```

. subst ha; simp
. rw [Finset.sum_eq_single (a * v)]
. simp [mul_inv_cancel_left ha]
. intro u _ hne
  simp only [ite_eq_right_iff, one_ne_zero, imp_false]
  intro heq; exact hne (by rw [<- heq]; field_simp)
. intro hv; exact absurd (Finset.mem_univ _) hv

-- S4: Simulator B1 -- pk hop (Lemma 5.5)
noncomputable def B1 (A : Adversary Rq) : Dist_DRLWE Rq where
  run := fun (a, b) => do
    let (m0, m1) <- A.choose (a, b)
    let r <- chi; let e1 <- chi; let e2 <- chi
    A.guess (a, b) (m0, m1) (a*r+e1, b*r+e2+m1)

lemma B1_real_eq_Game0 (A : Adversary Rq) :
  (do let a <- U; let s <- chi; let e <- chi;
    (B1 chi A).run (a, a*s+e)) = Game0 U chi A true := by
  simp only [B1, Game0]; ext x
  simp only [PMF.bind_apply]; congr 1; ext a; congr 1; ext s
  congr 1; ext e; simp [PMF.bind_apply]

lemma B1_unif_eq_Game1 (A : Adversary Rq) :
  (do let a <- U; let u <- U; (B1 chi A).run (a, u))
  = Game1 U chi A true := by
  simp only [B1, Game1]; ext x
  simp only [PMF.bind_apply]; congr 1; ext a; congr 1; ext u
  simp [PMF.bind_apply]

-- S5: Simulator B2 -- ciphertext hop (Lemma 5.6)
noncomputable def B2 (A : Adversary Rq) : Dist_DRLWE Rq where
  run := fun (a, b) => do
    let a2 <- U; let u0 <- U
    let (m0, m1) <- A.choose (a2, u0)
    let e2 <- chi
    A.guess (a2, u0) (m0, m1) (b, a*b+e2+m1)

-- S6: Game chain (Lemma 5.7)
lemma Game0_false_eq_Game2 (A : Adversary Rq) :
  (Game0 U chi A false).toFun true = (Game2 U A).toFun true := by
  have h1 : (Game0 U chi A false).toFun true
    = (Game1 U chi A false).toFun true :=
  drlwe_hardness (U := U) (chi := chi) { run := fun (a, b) => do
    let (m0, m1) <- A.choose (a, b)
    let r <- chi; let e1 <- chi; let e2 <- chi
    A.guess (a, b) (m0, m1) (a*r+e1, b*r+e2+m0) }
  have h2 : (Game1 U chi A false).toFun true
    = (Game2 U A).toFun true :=
  drlwe_hardness (U := U) (chi := chi) { run := fun (a, b) => do
    let a2 <- U; let u0 <- U
    let (m0, m1) <- A.choose (a2, u0); let e2 <- chi
    A.guess (a2, u0) (m0, m1) (b, a*b+e2+m0) }
  rw [h1, h2]

-- S7: Main theorem (Theorem 5.9)
/-- Adv_INDCPA(A) <= Adv_DRLWE(B1) + Adv_DRLWE(B2).
  Axiom count: 1. Sorry count: 0. -/
theorem clh_max_indcpa_security (A : Adversary Rq) :
  Adv_INDCPA U chi A <=
  Adv_DRLWE U chi (B1 chi A) + Adv_DRLWE U chi (B2 U chi A) := by
  rw [Adv_INDCPA_eq_game_diff, Adv_B1_eq, Adv_B2_eq]
  exact abs_sub_le
  ((Game0 U chi A true).toFun true)

```

```

      ((Game1 U chi A true).toFun true)
      ((Game2 U A).toFun true)

/-- Corollary 5.10: Adv_INDCPA = 0 under D-RLWE hardness. -/
corollary clh_max_indcpa_zero (A : Adversary Rq) :
  Adv_INDCPA U chi A = 0 := by
  have h := clh_max_indcpa_security U chi A
  rw [Adv_DRLWE_eq_zero, Adv_DRLWE_eq_zero] at h
  linarith [abs_nonneg (Adv_INDCPA U chi A)]
end ClhMax

```